

# Swin Transformer

Team #9

Sujy Han

Jaehyup Seong

Al Jamil, Mouhammed

2022.05.27

# • Index

## ➤ CNN and ViT

- CNN architecture
- Inductive bias
- Vision Transformer

## ➤ DeiT

- DeiT and CNN
- What is DeiT?
- DeiT and other Transformers

## ➤ Swin Transformers

- Problem statements
- Proposed methods
- Experiments
- Adaptation

# • Index

## ➤ CNN and ViT

- CNN architecture
- Inductive bias
- Vision Transformer


## ➤ DeiT

- DeiT and CNN
- What is DeiT?
- DeiT and other Transformers

## ➤ Swin Transformers

- Problem statements
- Proposed methods
- Experiments
- Adaptation

# • CNN Architecture



- The convolution operation derives the **translation equivariance**
- Pooling layer and Softmax function derive the **translation invariance**
- We call these 2 properties the **inductive bias** of CNN architecture

# • Inductive bias of CNN architecture


## ❖ Transition Invariance

- Although the location of the target is changed, the output of CNN architecture remains



## ❖ Transition Equivariance

- Although the location of the target is changed, the parameter of kernel is same



# • Inductive bias of CNN architecture

## ❖ Strong aspects of pre-defined properties

- Well predict most unseen images with pre-defined prediction
- Good performance even with the low number of training data




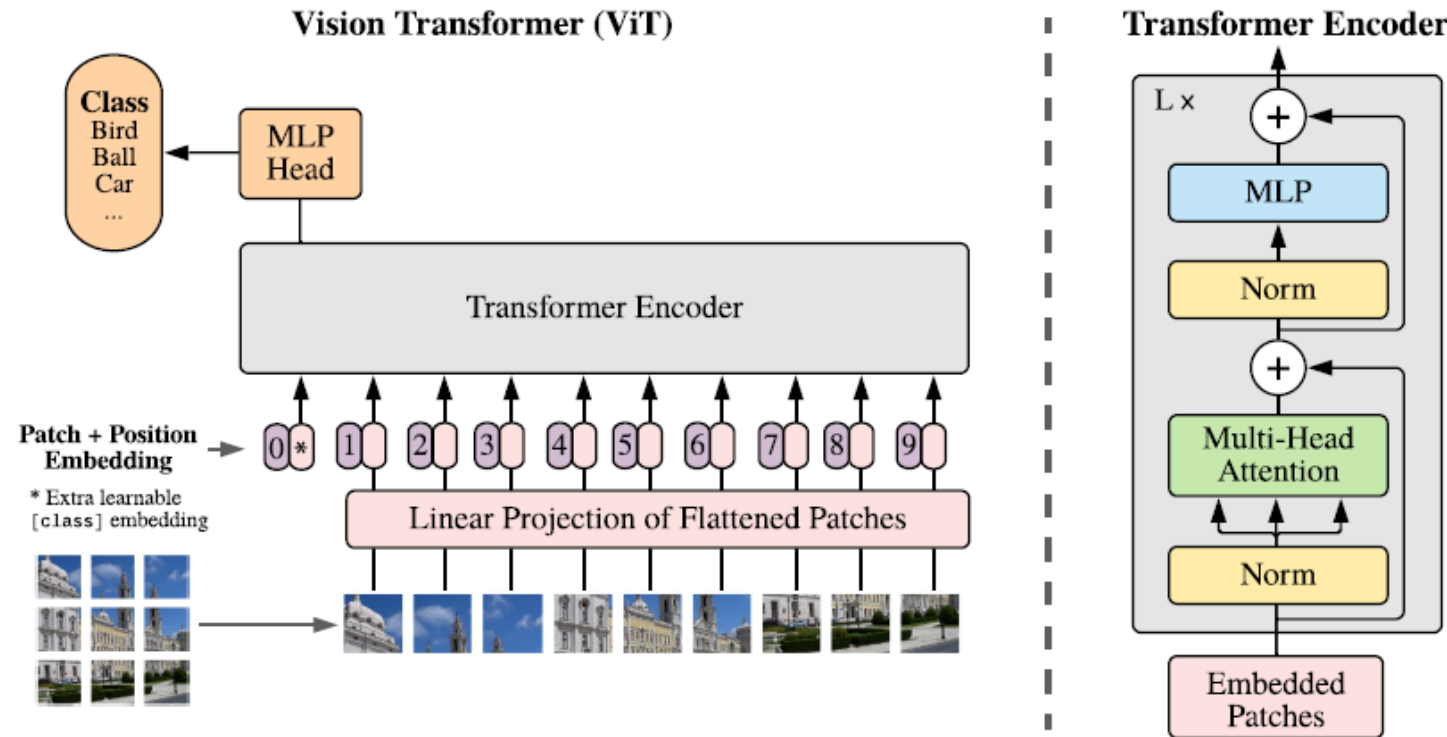
Figure credit: <https://github.com/aleju/imgaug>

## ❖ Weak aspects of pre-defined properties

- Loss of spatial features due to translation invariance
  - This can be overcome by data augmentation, but it demands a lot of cost overhead
- Poorly predict some images, out of the pre-defined induction (Poorly train about global information)



# • Vision Transformer (ViT)



- Idea from the transformer, which used for NLP
- ViT has no convolution operation and the pooling layer
- Thus, ViT has no **inductive bias** of translation equivariance and translation invariance

# • Vision Transformer (ViT)

## How can we deal without inductive bias?

### ❖ Huge amount of training data (e.g., JFT-300M)

- To predict the image without the pre-define induction, generalized properties of the image are required
- Generalized properties of images can be derived with a huge amount of data

### ❖ Without convolution operation and pooling layer,

*Attention is all you need*





# • Vision Transformer (ViT)

Image:  $\mathbf{x} \in \mathbb{R}^{H \times W \times C}$

$$N = HW/P^2$$

*1. To make flatten patches*

Patch:  $\mathbf{x}_p \in \mathbb{R}^{N \times (P^2 \cdot C)}$

$$\mathbf{E} \in \mathbb{R}^{(P^2 \cdot C) \times D}$$

*2. Linear projection*

Linearly projected patch:  $\mathbf{x}_p \mathbf{E} \in \mathbb{R}^{N \times D}$

$$\mathbf{x}_{class} \in \mathbb{R}^D$$


*3. Add the class token*

Add Class token:  $\mathbf{z}_0 = [\mathbf{x}_{class}; \mathbf{x}_p^1 \mathbf{E}; \mathbf{x}_p^2 \mathbf{E}; \dots; \mathbf{x}_p^N \mathbf{E}]$

$$\mathbf{E}_{pos} \in \mathbb{R}^{(N+1) \times D}$$

*4. Add the positioning vector*

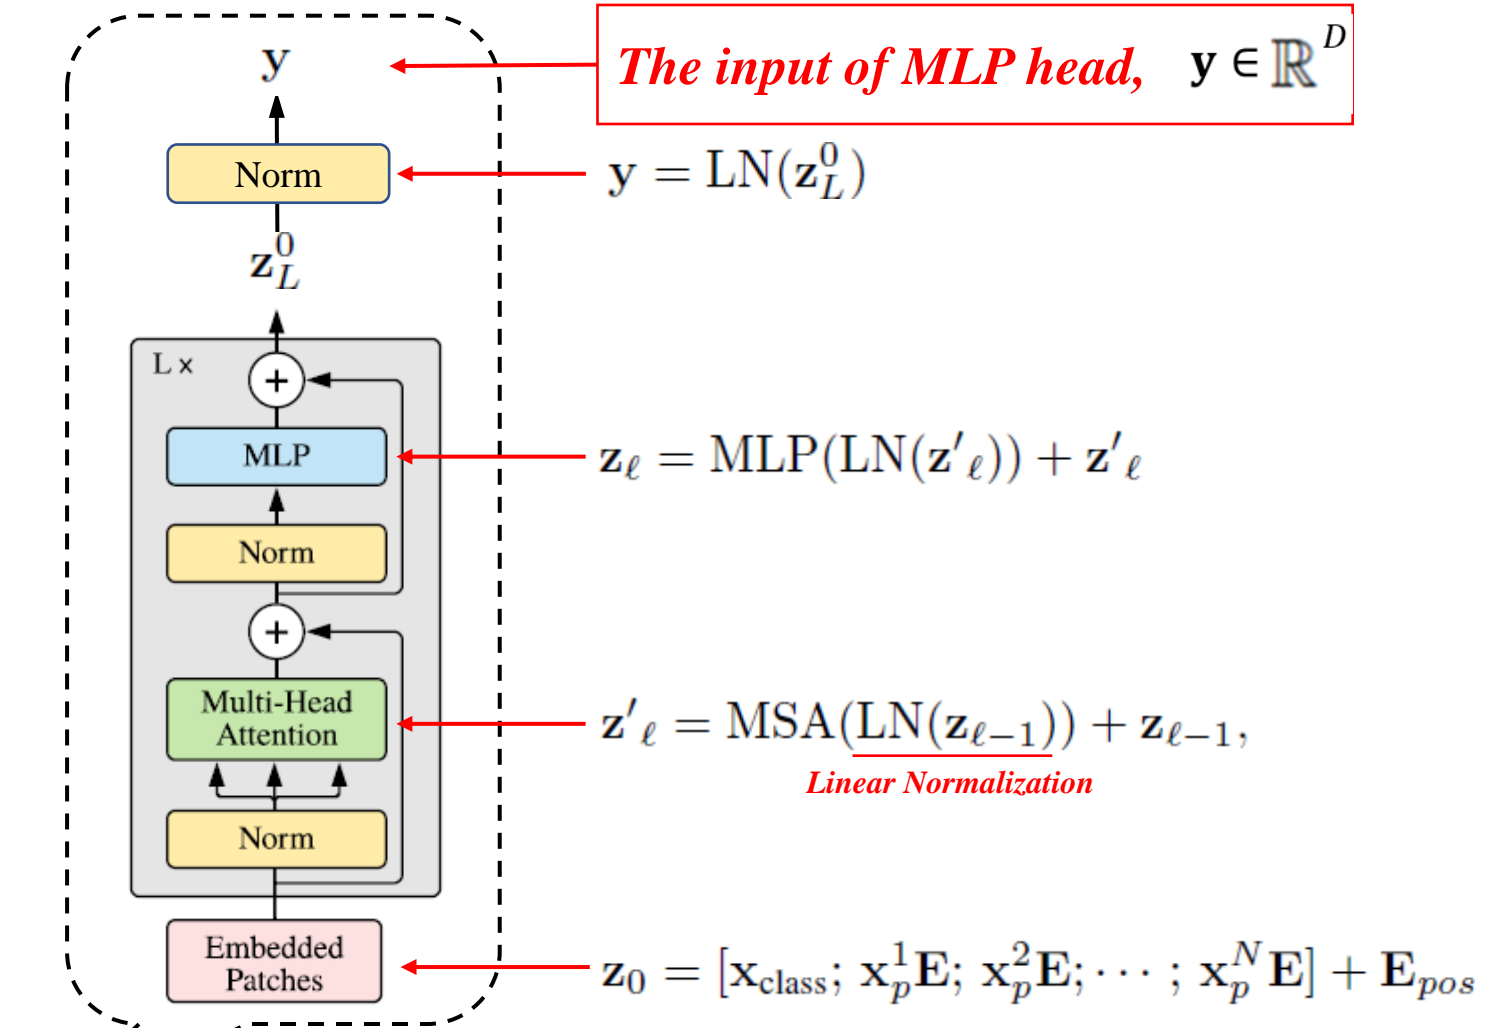
Embed positioning vector:  $\mathbf{z}_0 = [\mathbf{x}_{class}; \mathbf{x}_p^1 \mathbf{E}; \mathbf{x}_p^2 \mathbf{E}; \dots; \mathbf{x}_p^N \mathbf{E}] + \mathbf{E}_{pos}$



$$E_{(pos, 2i)} = \sin(pos/10000^{2i/d_{model}})$$

$$E_{(pos, 2i+1)} = \cos(pos/10000^{2i/d_{model}})$$

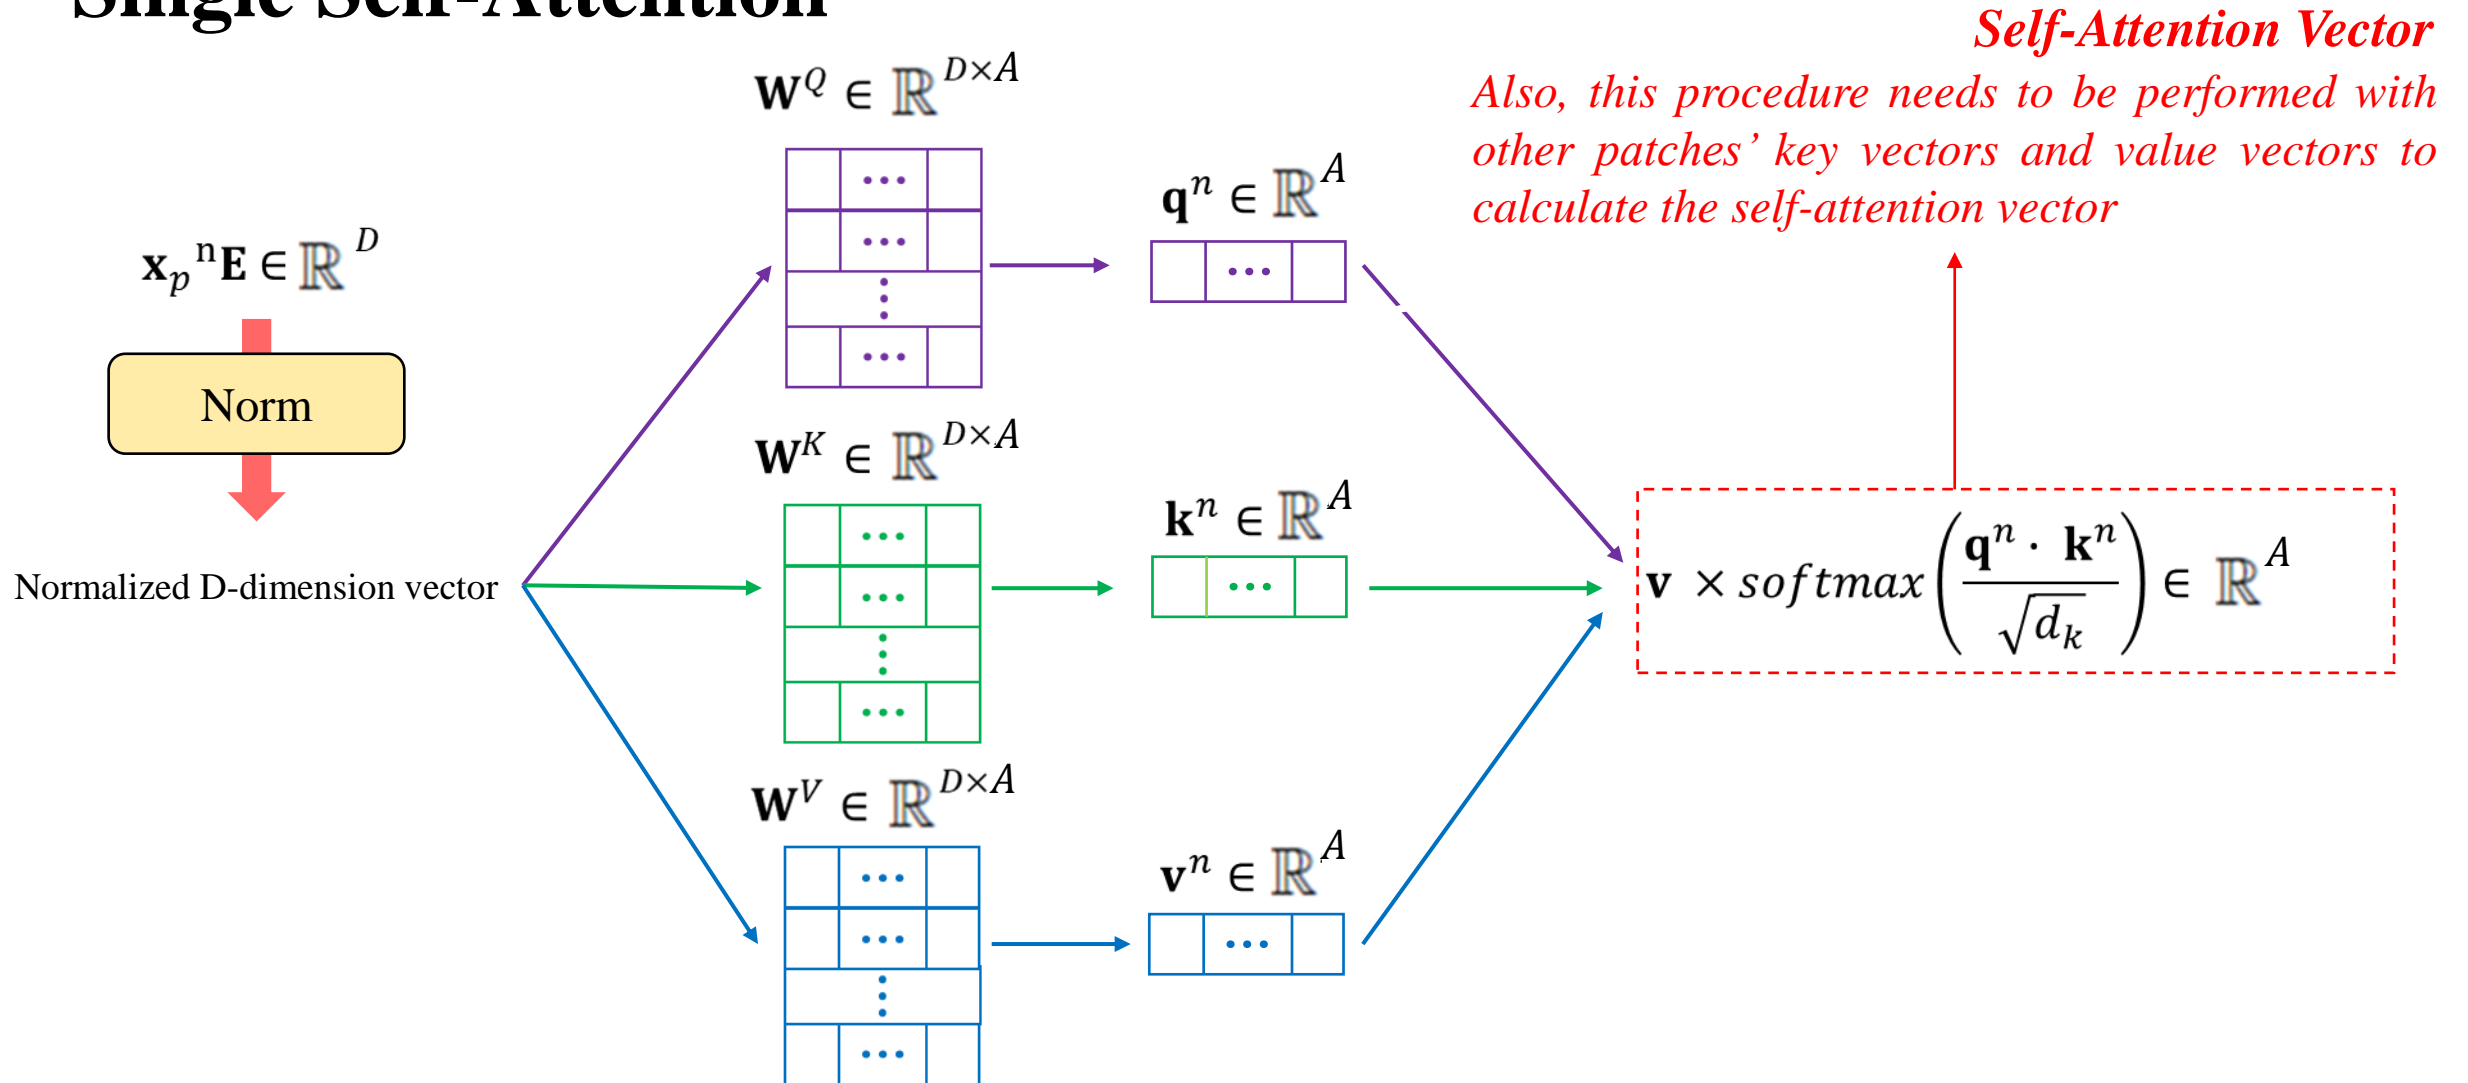
# • Vision Transformer (ViT)



Transformer Encoder

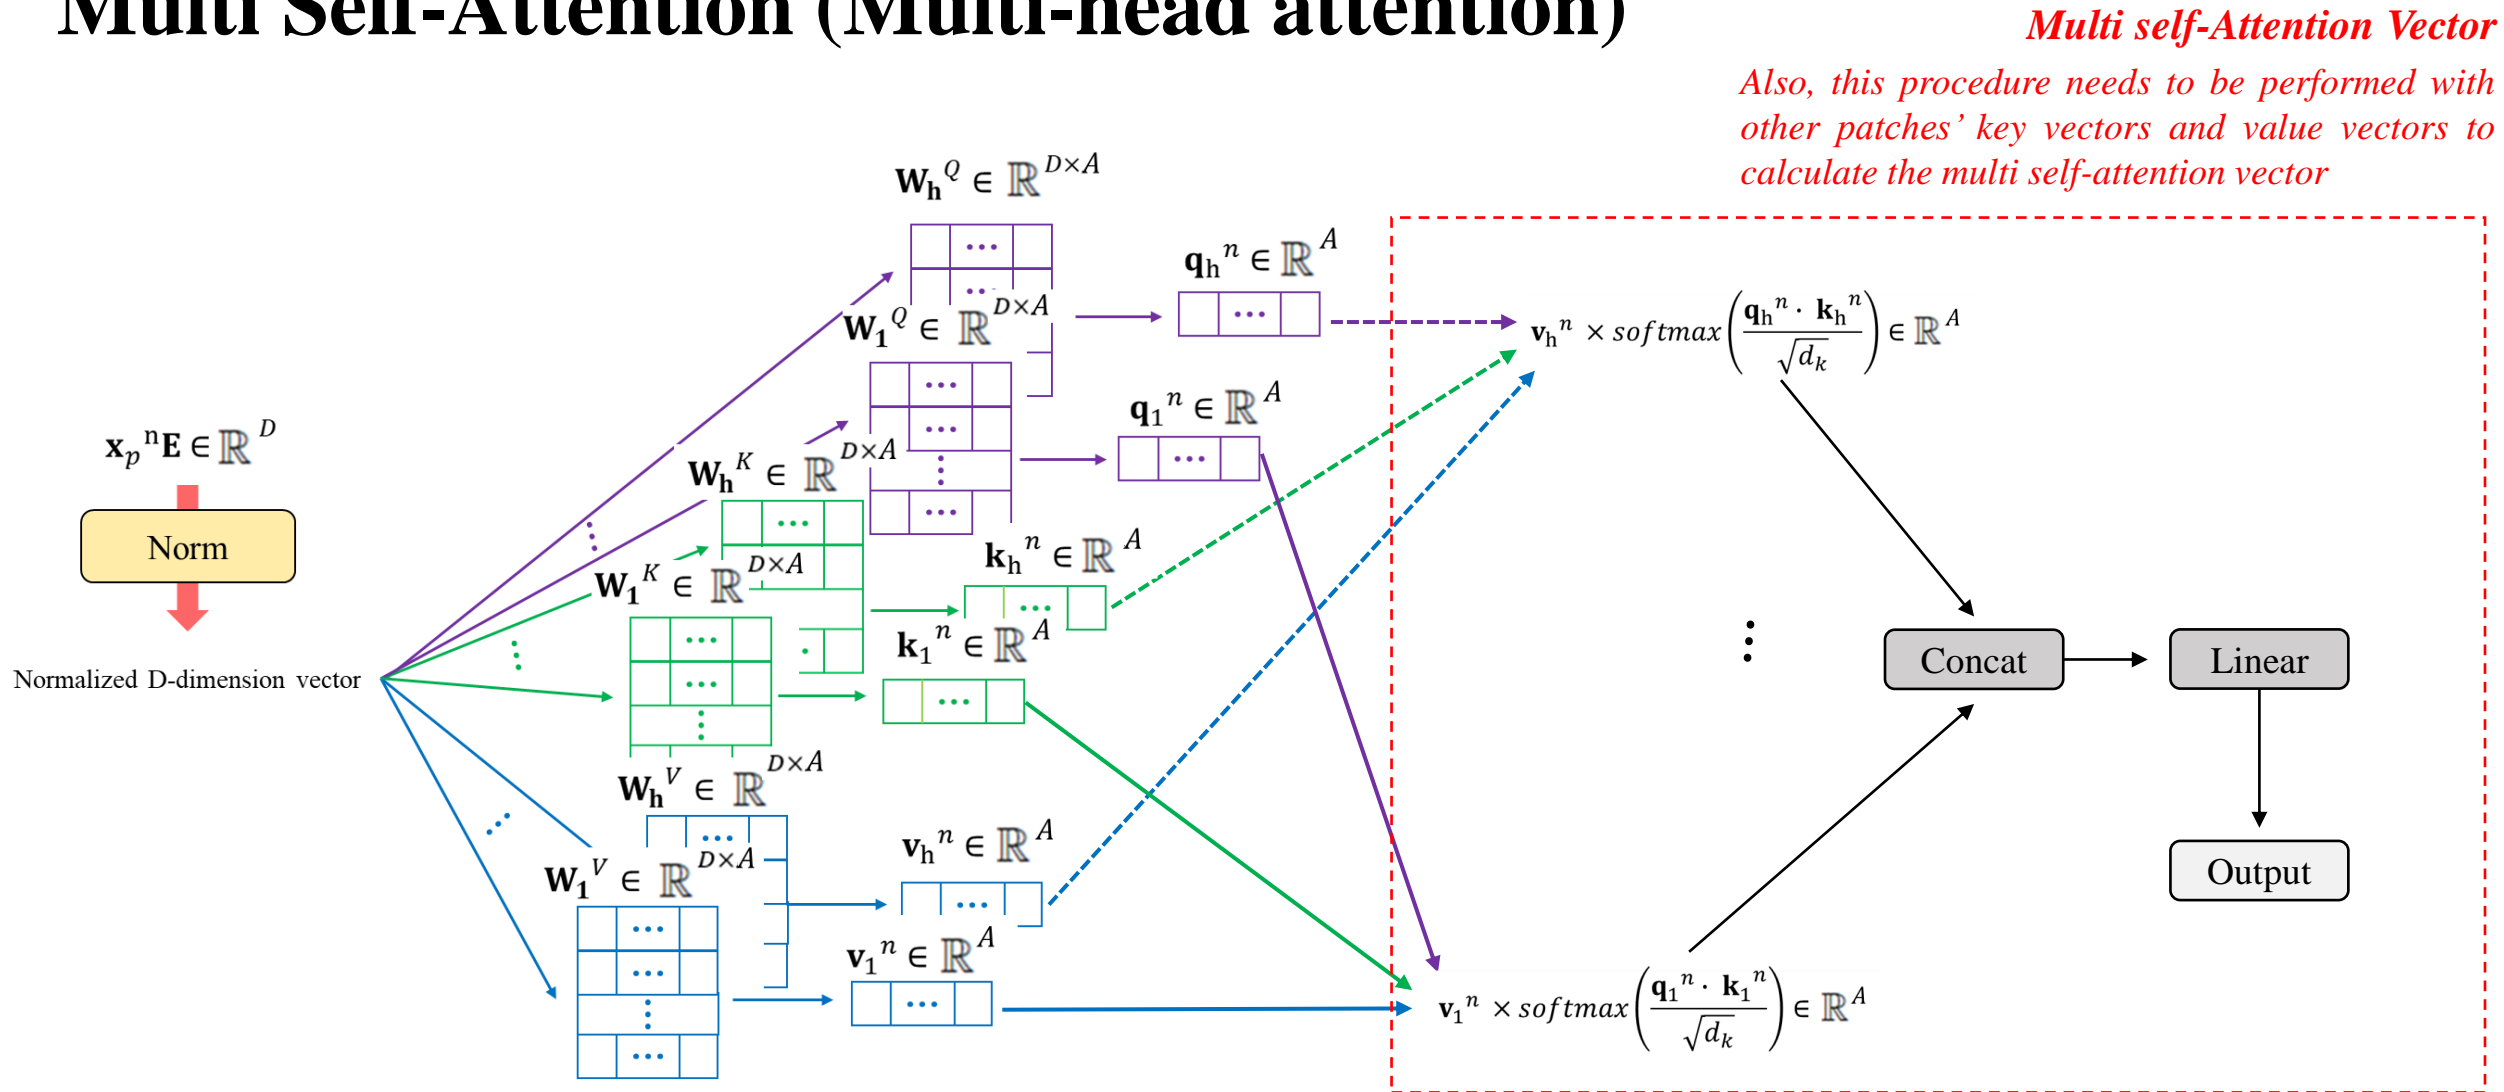
# • Vision Transformer (ViT)

## Single Self-Attention



# • Vision Transformer (ViT)

## Multi Self-Attention (Multi-head attention)



# • Vision Transformer (ViT)



## Pre-trained with JFT-300M dataset

|                    | Ours-JFT<br>(ViT-H/14)  | Ours-JFT<br>(ViT-L/16)  | Ours-I21k<br>(ViT-L/16) | BiT-L<br>(ResNet152x4) | Noisy Student<br>(EfficientNet-L2) |
|--------------------|-------------------------|-------------------------|-------------------------|------------------------|------------------------------------|
| ImageNet           | <b>88.55</b> $\pm 0.04$ | 87.76 $\pm 0.03$        | 85.30 $\pm 0.02$        | 87.54 $\pm 0.02$       | 88.4/88.5*                         |
| ImageNet ReaL      | <b>90.72</b> $\pm 0.05$ | 90.54 $\pm 0.03$        | 88.62 $\pm 0.05$        | 90.54                  | 90.55                              |
| CIFAR-10           | <b>99.50</b> $\pm 0.06$ | 99.42 $\pm 0.03$        | 99.15 $\pm 0.03$        | 99.37 $\pm 0.06$       | —                                  |
| CIFAR-100          | <b>94.55</b> $\pm 0.04$ | 93.90 $\pm 0.05$        | 93.25 $\pm 0.05$        | 93.51 $\pm 0.08$       | —                                  |
| Oxford-IIIT Pets   | <b>97.56</b> $\pm 0.03$ | 97.32 $\pm 0.11$        | 94.67 $\pm 0.15$        | 96.62 $\pm 0.23$       | —                                  |
| Oxford Flowers-102 | 99.68 $\pm 0.02$        | <b>99.74</b> $\pm 0.00$ | 99.61 $\pm 0.02$        | 99.63 $\pm 0.03$       | —                                  |
| VTAB (19 tasks)    | <b>77.63</b> $\pm 0.23$ | 76.28 $\pm 0.46$        | 72.72 $\pm 0.21$        | 76.29 $\pm 1.70$       | —                                  |
| TPUv3-core-days    | 2.5k                    | 0.68k                   | 0.23k                   | 9.9k                   | 12.3k                              |

- Although the number of patches increases when the size of the image gets bigger, the amount of GPU memory does not change much.
- Therefore, ViT can use the large batch size, which enables to use of high GPU utilization.
- High GPU utilization accelerates the training speed.

# • Vision Transformer (ViT)


## Importance of the size of pre-trained data



- Large ViT models perform worse than BiT ResNets when pre-trained on smaller datasets.
- Large ViT models perform better than ResNets when pre-trained by larger datasets.

# • Vision Transformer (ViT)

## Hybrid ViT



- Patches can be substituted for features of the image, which are formed from the feature maps of the CNN
- The input sequence is obtained by simply flattening the spatial dimension of the feature map
- Highly performs than casual ViT, when training cost is limited

# • Index

## ➤ CNN and ViT

- CNN architecture
- Inductive bias
- Vision Transformer

## ➤ DeiT

- DeiT and CNN
- What is DeiT?
- DeiT and other Transformers

## ➤ Swin Transformers

- Problem statements
- Proposed methods
- Experiments
- Adaptation



# DeiT



- **DeiT and ConvNets**


- is a class of artificial neural network (ANN), most commonly applied to analyze visual imagery
- Convolutional neural networks have been the main design paradigm for image understanding tasks
- They can be found at the core of everything from Facebook's photo tagging to self-driving cars.

- **What is DeiT?**

- DeiT: Data efficient image transformers
- DeiTs are a special kind of convolution free Transformers which involve s a teacher-student strategy that relies on a distillation token ensuring that the student learns from the teacher through attention.
- In DeiT new distillation procedure based on distillation tokens is implemented

- **DeiT**

## DeiT Architecture



- **DeiT**

## **Difference between DeiT and other Transformers**

- Normal vision transformer (ViT) needs a large volume of curated data in order to be effective.
- While using the DeiT approach the transformers can be trained on a single computer in less than 3 days (53 hours of pre-training, and optionally 20 hours of fine-tuning) with a top-1 accuracy of 83.1% on ImageNet with no external data (according to the result of the research paper experiment)

# • Index

## ➤ CNN and ViT

- CNN architecture
- Inductive bias
- Vision Transformer


## ➤ DeiT

- DeiT and CNN
- What is DeiT?
- DeiT and other Transformers

## ➤ Swin Transformers

- Problem statements
- Proposed methods
- Experiments
- Adaptation

- **Swin Transformers**



# • Swin Transformers

## What is Swin Transformers

- Swin Transformers are basically Hierarchical Vision Transformer using Shifted Windows.

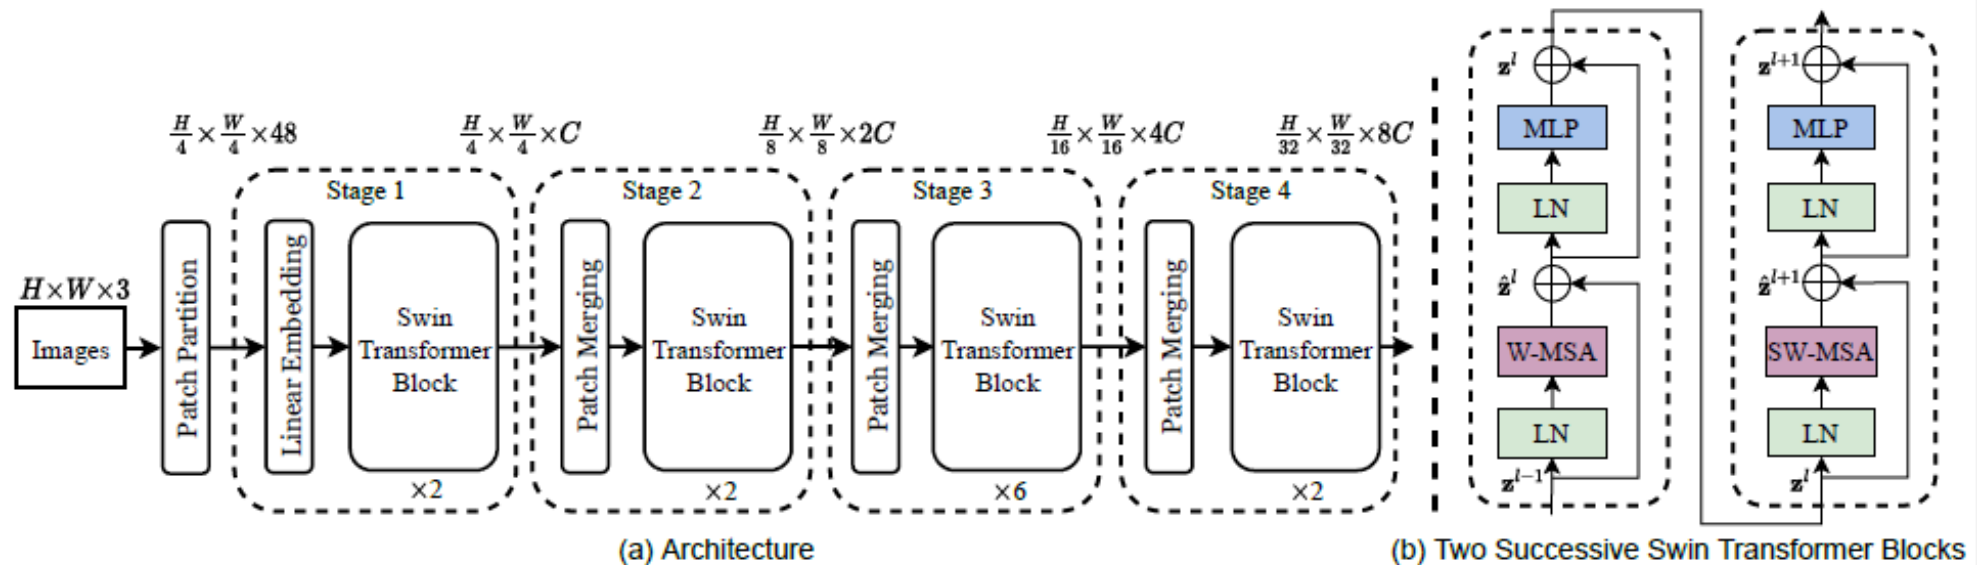


Figure 3. (a) The architecture of a Swin Transformer (Swin-T); (b) two successive Swin Transformer Blocks (notation presented with Eq. (3)). W-MSA and SW-MSA are multi-head self attention modules with regular and shifted windowing configurations, respectively.



- **Swin Transformers**

## **Problem Statement**

- Computer vision has long been dominated by convolutional neural networks (CNNs)
- Due to tremendous success in the NLP (Neuro-linguistic programming), it was adapted to computer vision such as image classification and joint vision-language modeling.
- **Scaling Problem:** Visual elements can vary substantially in scale, unlike the word tokens.
- The computational complexity of its self-attention is quadratic to image size.

# • Swin Transformers

## Overall Architecture

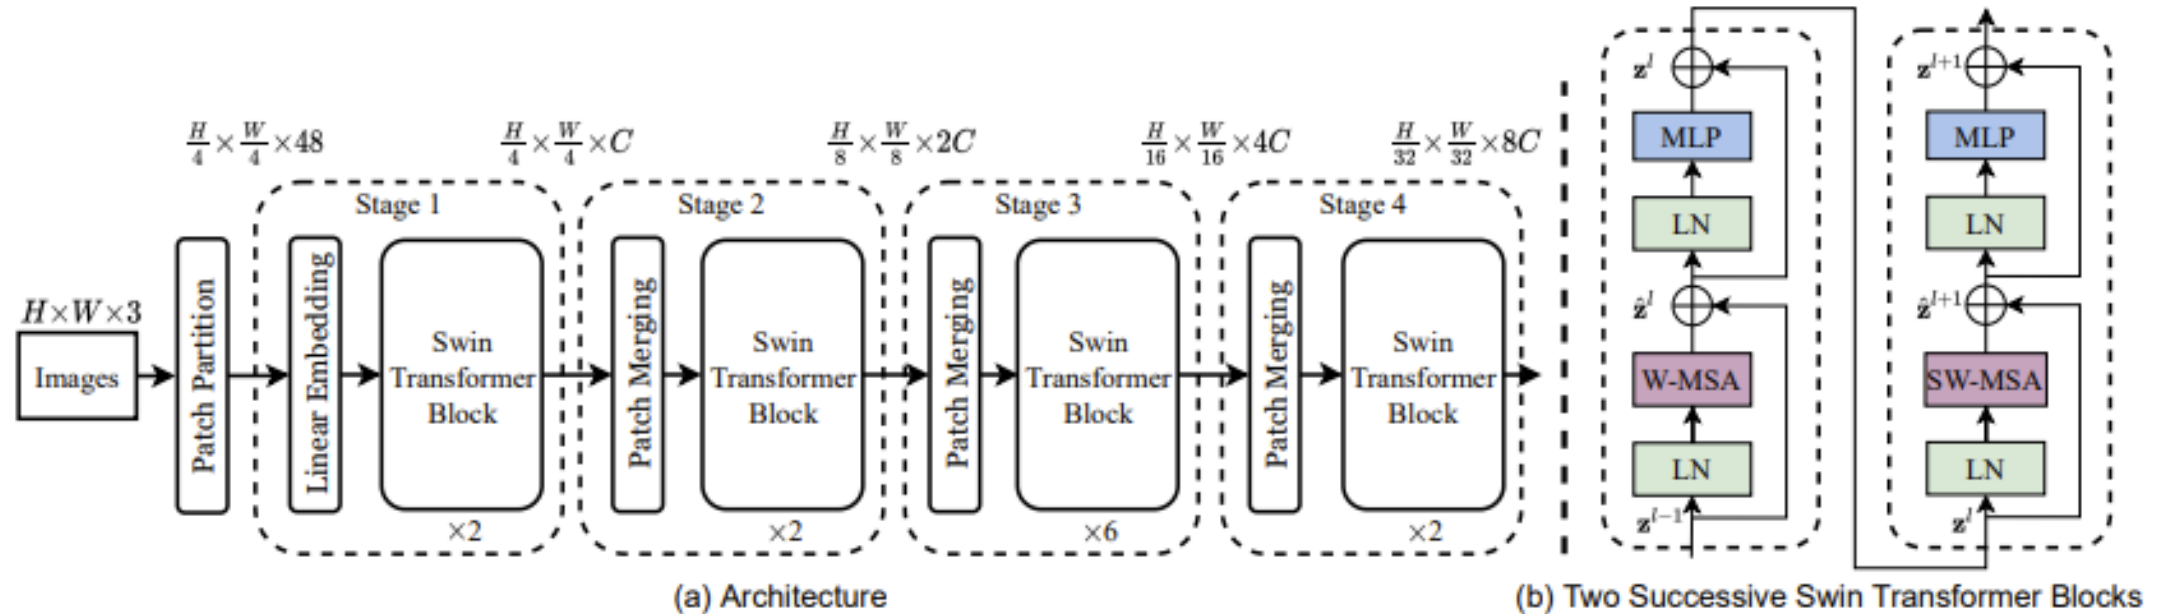


Figure 3. (a) The architecture of a Swin Transformer (Swin-T); (b) two successive Swin Transformer Blocks (notation presented with Eq. (3)). W-MSA and SW-MSA are multi-head self attention modules with regular and shifted windowing configurations, respectively.

- Figure 3 illustrates the tiny version of Swin Transformer.
- It consists of Patch Partition, Patch Merging and Swin Transformer Block.

# • Swin Transformers

## Overall Architecture : Patch Partition




Figure 3. (a) The architecture of a Swin Transformer (Swin-T); (b) two successive Swin Transformer Blocks (notation presented with Eq. (3)). W-MSA and SW-MSA are multi-head self attention modules with regular and shifted windowing configurations, respectively.

- An input RGB image ( $H \times W \times 3$ ) is split into **non-overlapping patches**.
- Each patch is treated as a **token**.

# • Swin Transformers

## Overall Architecture : Stage 1

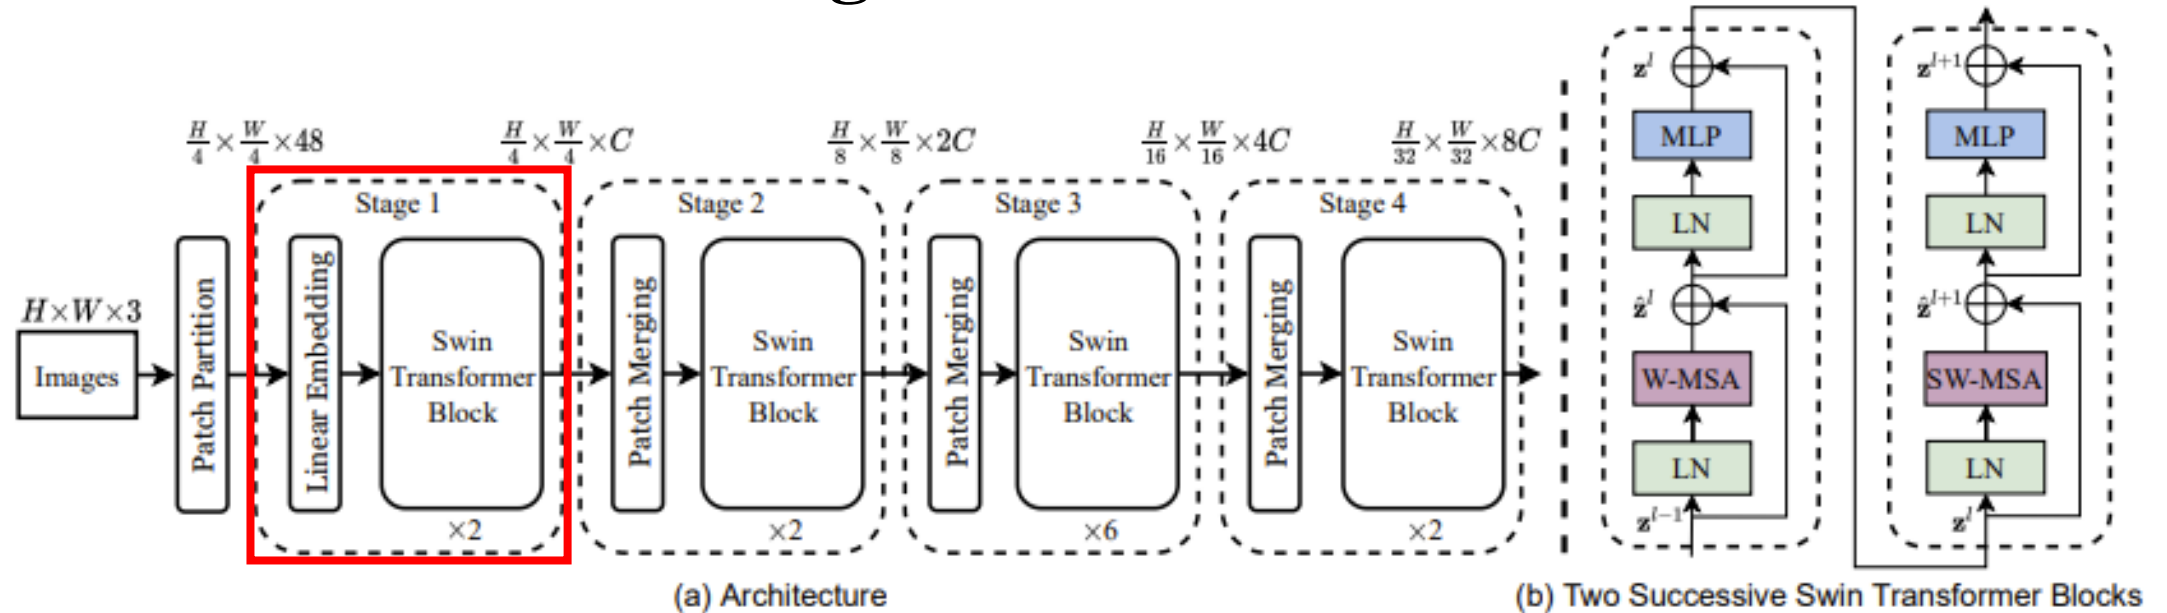


Figure 3. (a) The architecture of a Swin Transformer (Swin-T); (b) two successive Swin Transformer Blocks (notation presented with Eq. (3)). W-MSA and SW-MSA are multi-head self attention modules with regular and shifted windowing configurations, respectively.

- **Linear embedding layer** is applied on an input tensor.
- Instead of general MSA (Multi-head Self Attention), **W-MSA** and **SW-MSA** are used.

# • Swin Transformers

## Overall Architecture : Stage 2~4




Figure 3. (a) The architecture of a Swin Transformer (Swin-T); (b) two successive Swin Transformer Blocks (notation presented with Eq. (3)). W-MSA and SW-MSA are multi-head self attention modules with regular and shifted windowing configurations, respectively.

- The number of tokens is reduced by patch merging layers.
- The first patch merging layer concatenates the features of each group of **2x2 neighboring patches**.
- **Swin Transformer blocks** are applied after the feature transformation.
- These stages produce a **hierarchical representation**, which helps comparative abilities in computer vision.

- # Swin Transformers

## Overall Architecture : Swin Transformer Block

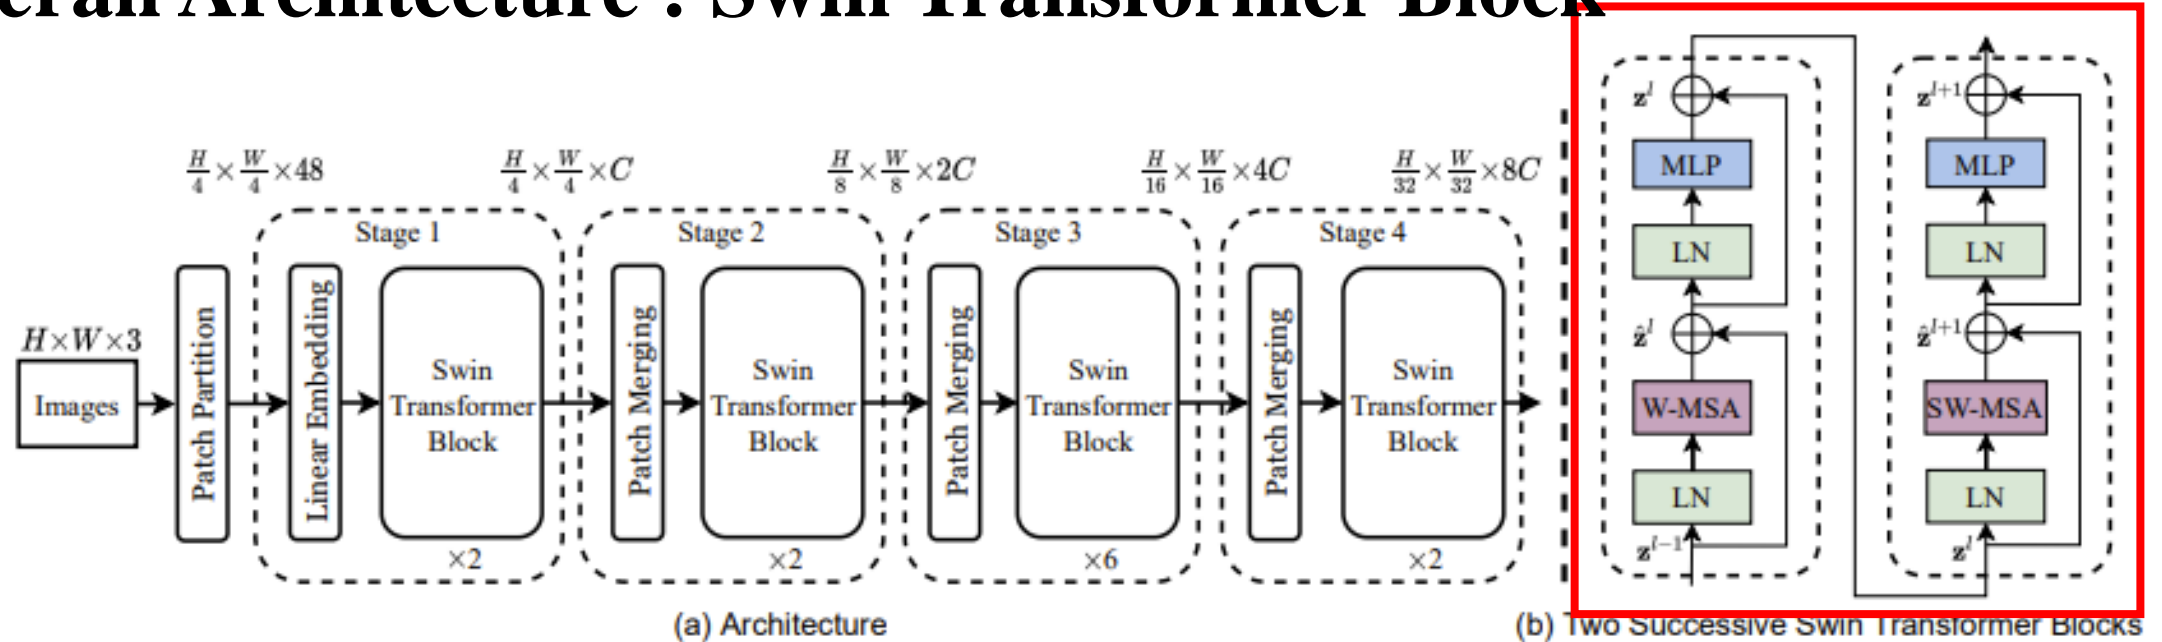


Figure 3. (a) The architecture of a Swin Transformer (Swin-T); (b) two successive Swin Transformer Blocks (notation presented with Eq. (3)). W-MSA and SW-MSA are multi-head self attention modules with regular and shifted windowing configurations, respectively.

- Swin Transformer is built by **replacing the standard MSA module** in a Transformer block **by a module based on shifted windows**.
- A Swin Transformer block consists of a **shifted window based MSA module**, followed by a 2-layer MLP with GELU non-linearity in between.



# • Swin Transformers

## Shifted Window based Self-Attention

$$\Omega(\text{MSA}) = 4hwC^2 + 2(hw)^2C, \quad (1)$$

$$\Omega(\text{W-MSA}) = 4hwC^2 + 2M^2hwC, \quad (2)$$




Figure 2. An illustration of the *shifted window* approach for computing self-attention in the proposed Swin Transformer architecture. In layer  $l$  (left), a regular window partitioning scheme is adopted, and self-attention is computed within each window. In the next layer  $l + 1$  (right), the window partitioning is shifted, resulting in new windows. The self-attention computation in the new windows crosses the boundaries of the previous windows in layer  $l$ , providing connections among them.

- **Global self-attention** computes the relationships between a token and all other tokens, which leads to **quadratic complexity** w.r.t. the number of tokens.
- **Self-attention within local-windows** is proposed for efficient modeling.
- (1) is quadratic to patch number  $hw$  while (2) is **linear** when  $M$  is fixed.
- **Shifted window partitioning** approach is to introduce **cross-window connections**.

# • Swin Transformers

## Efficient batch computation for shifted configuration

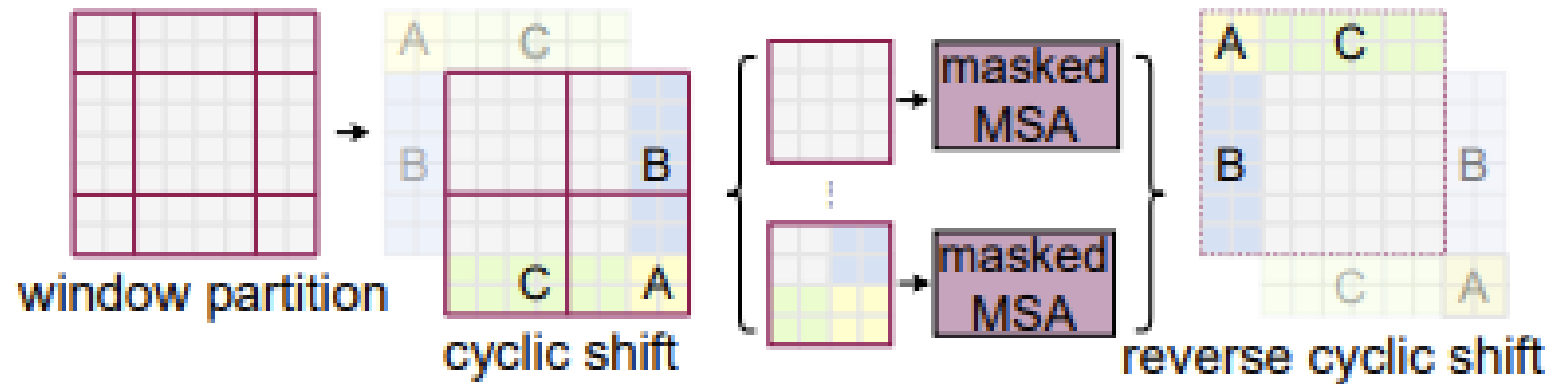


Figure 4. Illustration of an efficient batch computation approach for self-attention in shifted window partitioning.

- In the shifted configuration, some windows will be **smaller than  $M \times M$**
- For more efficient computation, **cyclic-shifting** toward the top-left direction is proposed.
- A **masking mechanism** is employed to limit self-attention computation to within each sub-window.



# • Swin Transformers

## Architecture Variants

- Swin-T:  $C = 96$ , layer numbers =  $\{2, 2, 6, 2\}$
- Swin-S:  $C = 96$ , layer numbers =  $\{2, 2, 18, 2\}$
- Swin-B:  $C = 128$ , layer numbers =  $\{2, 2, 18, 2\}$
- Swin-L:  $C = 192$ , layer numbers =  $\{2, 2, 18, 2\}$

- $C$  is the channel number of the hidden layers in the first stage.
- Swin-B is the base model, which has the similar model size and computation complexity to ViT-B and DeiT-B.
- Swin-T, Swin-S, Swin-L are versions of about 0.25x, 0.5x and 2x the model size and computational complexity, respectively.

- **Swin Transformers**

## Image Classification on ImageNet-1K / 22K

| (a) Regular ImageNet-1K trained models |                  |         |        |                        |                     |
|--|------------------|---------|--------|------------------------|---------------------|
| method                                 | image size       | #param. | FLOPs  | throughput (image / s) | ImageNet top-1 acc. |
| RegNetY-4G [44]                        | 224 <sup>2</sup> | 21M     | 4.0G   | 1156.7                 | 80.0                |
| RegNetY-8G [44]                        | 224 <sup>2</sup> | 39M     | 8.0G   | 591.6                  | 81.7                |
| RegNetY-16G [44]                       | 224 <sup>2</sup> | 84M     | 16.0G  | 334.7                  | 82.9                |
| ViT-B/16 [19]                          | 384 <sup>2</sup> | 86M     | 55.4G  | 85.9                   | 77.9                |
| ViT-L/16 [19]                          | 384 <sup>2</sup> | 307M    | 190.7G | 27.3                   | 76.5                |
| DeiT-S [57]                            | 224 <sup>2</sup> | 22M     | 4.6G   | 940.4                  | 79.8                |
| DeiT-B [57]                            | 224 <sup>2</sup> | 86M     | 17.5G  | 292.3                  | 81.8                |
| DeiT-B [57]                            | 384 <sup>2</sup> | 86M     | 55.4G  | 85.9                   | 83.1                |
| Swin-T                                 | 224 <sup>2</sup> | 29M     | 4.5G   | 755.2                  | 81.3                |
| Swin-S                                 | 224 <sup>2</sup> | 50M     | 8.7G   | 436.9                  | 83.0                |
| Swin-B                                 | 224 <sup>2</sup> | 88M     | 15.4G  | 278.1                  | 83.5                |
| Swin-B                                 | 384 <sup>2</sup> | 88M     | 47.0G  | 84.7                   | 84.5                |

| (b) ImageNet-22K pre-trained models |                  |         |        |                        |                     |
|-------------------------------------|------------------|---------|--------|------------------------|---------------------|
| method                              | image size       | #param. | FLOPs  | throughput (image / s) | ImageNet top-1 acc. |
| R-101x3 [34]                        | 384 <sup>2</sup> | 388M    | 204.6G | -                      | 84.4                |
| R-152x4 [34]                        | 480 <sup>2</sup> | 937M    | 840.5G | -                      | 85.4                |
| ViT-B/16 [19]                       | 384 <sup>2</sup> | 86M     | 55.4G  | 85.9                   | 84.0                |
| ViT-L/16 [19]                       | 384 <sup>2</sup> | 307M    | 190.7G | 27.3                   | 85.2                |
| Swin-B                              | 224 <sup>2</sup> | 88M     | 15.4G  | 278.1                  | 85.2                |
| Swin-B                              | 384 <sup>2</sup> | 88M     | 47.0G  | 84.7                   | 86.4                |
| Swin-L                              | 384 <sup>2</sup> | 197M    | 103.9G | 42.1                   | 87.3                |

Table 1. Comparison of different backbones on ImageNet-1K classification. Throughput is measured using the GitHub repository of [62] and a V100 GPU, following [57].

- # Swin Transformers

## Object Detection on COCO

| (a) Various frameworks                      |                   |                                 |                                 |                                 |                                  |                                  |                |
|---|-------------------|---------------------------------|---------------------------------|---------------------------------|----------------------------------|----------------------------------|----------------|
| Method                                      | Backbone          | AP <sup>box</sup>               | AP <sup>box</sup> <sub>50</sub> | AP <sup>box</sup> <sub>75</sub> | #param.                          | FLOPs                            | FPS            |
| Cascade                                     | R-50              | 46.3                            | 64.3                            | 50.5                            | 82M                              | 739G                             | 18.0           |
| Mask R-CNN                                  | Swin-T            | <b>50.5</b>                     | <b>69.3</b>                     | <b>54.9</b>                     | 86M                              | 745G                             | 15.3           |
| ATSS  | R-50              | 43.5                            | 61.9                            | 47.0                            | 32M                              | 205G                             | 28.3           |
|   | Swin-T            | <b>47.2</b>                     | <b>66.5</b>                     | <b>51.3</b>                     | 36M                              | 215G                             | 22.3           |
| RepPointsV2                                 | R-50              | 46.5                            | 64.6                            | 50.3                            | 42M                              | 274G                             | 13.6           |
|   | Swin-T            | <b>50.0</b>                     | <b>68.5</b>                     | <b>54.2</b>                     | 45M                              | 283G                             | 12.0           |
| Sparse R-CNN                                | R-50              | 44.5                            | 63.4                            | 48.2                            | 106M                             | 166G                             | 21.0           |
|   | Swin-T            | <b>47.9</b>                     | <b>67.3</b>                     | <b>52.3</b>                     | 110M                             | 172G                             | 18.4           |
| (b) Various backbones w. Cascade Mask R-CNN |                   |                                 |                                 |                                 |                                  |                                  |                |
|   | AP <sup>box</sup> | AP <sup>box</sup> <sub>50</sub> | AP <sup>box</sup> <sub>75</sub> | AP <sup>mask</sup>              | AP <sup>mask</sup> <sub>50</sub> | AP <sup>mask</sup> <sub>75</sub> | #paramFLOPsFPS |
| DeiT-S <sup>†</sup>                         | 48.0              | 67.2                            | 51.7                            | 41.4                            | 64.2                             | 44.3                             | 80M 889G 10.4  |
| R50   | 46.3              | 64.3                            | 50.5                            | 40.1                            | 61.7                             | 43.4                             | 82M 739G 18.0  |
| Swin-T                                      | <b>50.5</b>       | <b>69.3</b>                     | <b>54.9</b>                     | <b>43.7</b>                     | <b>66.6</b>                      | <b>47.1</b>                      | 86M 745G 15.3  |
| X101-32                                     | 48.1              | 66.5                            | 52.4                            | 41.6                            | 63.9                             | 45.2                             | 101M 819G 12.8 |
| Swin-S                                      | <b>51.8</b>       | <b>70.4</b>                     | <b>56.3</b>                     | <b>44.7</b>                     | <b>67.9</b>                      | <b>48.5</b>                      | 107M 838G 12.0 |
| X101-64                                     | 48.3              | 66.4                            | 52.3                            | 41.7                            | 64.0                             | 45.1                             | 140M 972G 10.4 |
| Swin-B                                      | <b>51.9</b>       | <b>70.9</b>                     | <b>56.5</b>                     | <b>45.0</b>                     | <b>68.4</b>                      | <b>48.7</b>                      | 145M 982G 11.6 |

| (c) System-level Comparison |                   |                    |                   |                    |               |
|-----------------------------|-------------------|--------------------|-------------------|--------------------|---------------|
| Method                      | mini-val          |                    | test-dev          |                    | #param. FLOPs |
|                             | AP <sup>box</sup> | AP <sup>mask</sup> | AP <sup>box</sup> | AP <sup>mask</sup> |               |
| RepPointsV2* [12]           | -                 | -                  | 52.1              | -                  | - -           |
| GCNet* [7]                  | 51.8              | 44.7               | 52.3              | 45.4               | - 1041G       |
| RelationNet++* [13]         | -                 | -                  | 52.7              | -                  | - -           |
| DetectoRS* [42]             | -                 | -                  | 55.7              | 48.5               | - -           |
| YOLOv4 P7* [4]              | -                 | -                  | 55.8              | -                  | - -           |
| Copy-paste [23]             | 55.9              | 47.2               | 56.0              | 47.4               | 185M 1440G    |
| X101-64 (HTC++)             | 52.3              | 46.0               | -                 | -                  | 155M 1033G    |
| Swin-B (HTC++)              | 56.4              | 49.1               | -                 | -                  | 160M 1043G    |
| Swin-L (HTC++)              | 57.1              | 49.5               | 57.7              | 50.2               | 284M 1470G    |
| Swin-L (HTC++)*             | <b>58.0</b>       | <b>50.4</b>        | <b>58.7</b>       | <b>51.1</b>        | 284M -        |

Table 2. Results on COCO object detection and instance segmentation. <sup>†</sup>denotes that additional deconvolution layers are used to produce hierarchical feature maps. \* indicates multi-scale testing.

- # Swin Transformers

## Semantic Segmentation on ADE20K

| ADE20K        |                      | val         | test        | #param. | FLOPs | FPS  |
|---------------|----------------------|-------------|-------------|---------|-------|------|
| Method        | Backbone             | mIoU        | score       |         |       |      |
| DLab.v3+ [11] | ResNet-101           | 44.1        | -           | 63M     | 1021G | 16.0 |
| DNL [65]      | ResNet-101           | 46.0        | 56.2        | 69M     | 1249G | 14.8 |
| OCRNet [67]   | ResNet-101           | 45.3        | 56.0        | 56M     | 923G  | 19.3 |
| UperNet [63]  | ResNet-101           | 44.9        | -           | 86M     | 1029G | 20.1 |
| OCRNet [67]   | HRNet-w48            | 45.7        | -           | 71M     | 664G  | 12.5 |
| DLab.v3+ [11] | ResNeSt-101          | 46.9        | 55.1        | 66M     | 1051G | 11.9 |
| DLab.v3+ [11] | ResNeSt-200          | 48.4        | -           | 88M     | 1381G | 8.1  |
| SETR [73]     | T-Large <sup>‡</sup> | 50.3        | 61.7        | 308M    | -     | -    |
| UperNet       | DeiT-S <sup>†</sup>  | 44.0        | -           | 52M     | 1099G | 16.2 |
| UperNet       | Swin-T               | 46.1        | -           | 60M     | 945G  | 18.5 |
| UperNet       | Swin-S               | 49.3        | -           | 81M     | 1038G | 15.2 |
| UperNet       | Swin-B <sup>‡</sup>  | 51.6        | -           | 121M    | 1841G | 8.7  |
| UperNet       | Swin-L <sup>‡</sup>  | <b>53.5</b> | <b>62.8</b> | 234M    | 3230G | 6.2  |

Table 3. Results of semantic segmentation on the ADE20K val and test set. <sup>†</sup> indicates additional deconvolution layers are used to produce hierarchical feature maps. <sup>‡</sup> indicates that the model is pre-trained on ImageNet-22K.

# • Swin Transformers

## Contributions

- Swin Transformer presents a new vision Transformer.
- It produces a hierarchical feature representation.
- Its shifted window based self-attention has linear computational complexity with respect to input image size.

- **Swin Transformers**

## **SwinIR : Image Restoration Using Swin Transformer**

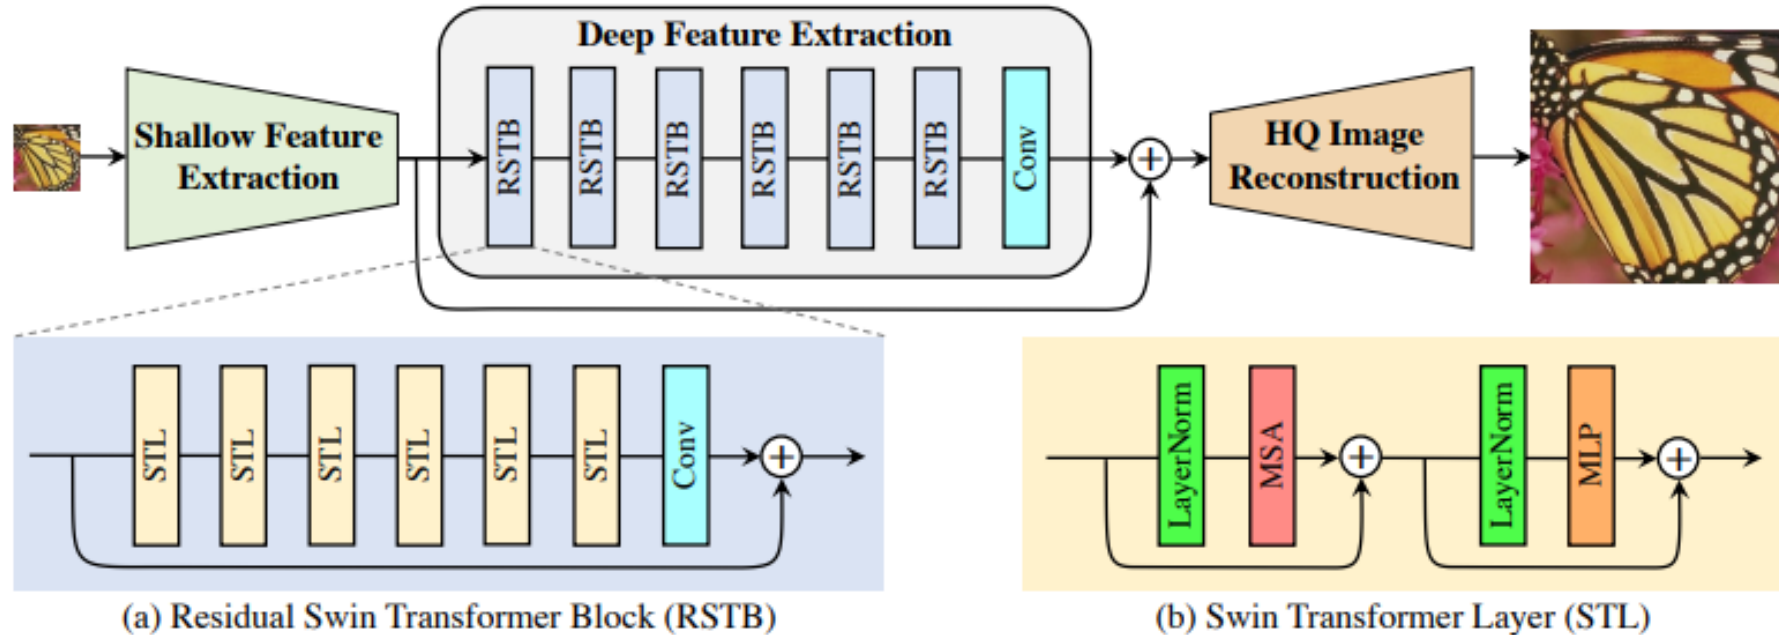


Figure 2: The architecture of the proposed SwinIR for image restoration.

- **An image restoration model based on Swin Transformer**
- SwinIR consists of three modules : shallow feature extraction, deep feature extraction and high-quality image reconstruction modules.



- Swin Transformers

## SwinIR : Image Restoration Using Swin Transformer




Figure 4: Visual comparison of bicubic image SR ( $\times 4$ ) methods. Compared images are derived from [9]. Best viewed by zooming.




Figure 5: Visual comparison of real-world image SR ( $\times 4$ ) methods on RealSRSet [89]. Compared images are derived from [89].



- # Swin Transformers

## SwinIR : Image Restoration Using Swin Transformer




Figure 6: Visual comparison of grayscale image denoising (noise level 50) methods on image “*Monarch*” from Set12 [90]. Compared images are derived from [88].



Figure 7: Visual comparison of color image denoising (noise level 50) methods on image “163085” from CBSD68 [59]. Compared images are derived from [88].